GrabOn

Vibe Coder Challenge

8 Project Briefs for Exceptional Builders

Format	8 independent project briefs — candidates choose one
Difficulty	Hard — each project requires 2.5 to 3 days to complete well
Tools Expected	Claude Code, MCP creation, REST API integration, frontend UI
Submission	GitHub repo + live demo + short Loom walkthrough
New Verticals	GrabCredit (fintech/lending) and GrabInsurance (embedded insurance)
Strategic Context	Rakuten, PayU, InMobi, Poonawalla Fincorp partnership frameworks
Evaluation	Technical depth + product thinking + demo quality + code quality

Context & Background

About GrabOn

GrabOn is India's #1 coupon and deals platform, processing 96M+ transactions per year across 3,500+ active merchants and 21,000+ total merchant relationships. With 40M subscribers and $4.8B in total GMV, GrabOn sits at a unique intersection: deep consumer trust in deal discovery, and deep merchant relationships across every major category in Indian e-commerce.

GrabOn's category footprint — Fashion & Beauty (24%), Travel (17%), Food (16%), Electronics (10%), Health (8%) — maps almost perfectly to the highest-value insurance and credit use cases. This is the foundational insight behind the two new verticals described in these briefs.

New Verticals: GrabCredit & GrabInsurance

GrabOn is building two new financial services verticals powered by its transaction data advantage and its strategic partnerships with Poonawalla Fincorp (an NBFC with lending license) and PayU (India's leading payment infrastructure).

GrabCredit will offer both consumer-facing BNPL products at checkout and merchant-facing working capital credit lines, underwritten using GrabOn's proprietary behavioral transaction data. GrabInsurance will offer contextual, embedded micro-insurance products served at the precise moment of deal redemption — travel insurance on a flight deal, electronics warranty on a gadget deal, health cover on a pharmacy deal.

Strategic Partnership Context

These project briefs are grounded in four active strategic partnerships, each representing a specific infrastructure or distribution layer:

Partner	Role	Relevance to These Projects
Rakuten	Affiliate & attribution infrastructure	Projects 03 (attribution platform) and 06 (deal distribution MCP)
Poonawalla Fincorp	NBFC lending license + embedded finance	Projects 01 (BNPL), 05 (wallet + FD), 07 (merchant underwriting)
InMobi / Glance	Mobile advertising + lock screen distribution	Project 04 (zero-input deal personalization for Glance)
PayU	Payment rails + LazyPay BNPL + checkout integration	Projects 01 (BNPL checkout), 05 (wallet spend flow)

What We Are Looking For in Candidates

These are not tutorial projects. A candidate who downloads a boilerplate and runs Claude on it without deep understanding will produce an obvious, shallow submission. We are looking for builders who:

- Understand the business problem deeply before writing a single line of code — the best submissions will show clear product thinking in their architecture decisions
- Use Claude Code and MCPs as force multipliers, not crutches — the code they write around the AI calls matters as much as the prompts
- Build for demo quality: the product must be impressive in a 10-minute live walkthrough to a non-technical audience (a Poonawalla Fincorp executive, a Rakuten account manager)
- Document their architecture, their API design decisions, and their tradeoffs clearly in a README — we read READMEs before running code

Project 01: GrabCredit BNPL Eligibility Engine

Build an AI-powered Buy Now, Pay Later credit scoring system at checkout, powered by GrabOn transaction history and PayU / Poonawalla Fincorp lending rails

[GrabCredit] [PayU / LazyPay] [Poonawalla Fincorp] [MCP Server] [Claude Agent]

Difficulty	Hard — 2.5 to 3 days
Vertical	GrabCredit (New Vertical)

The Challenge

GrabOn processes 96M+ transactions per year across 3,500+ active merchants. This mountain of behavioral data is GrabCredit's core asset — a uniquely rich signal for creditworthiness that traditional lenders don't have. This project challenges candidates to build the complete BNPL eligibility pipeline: from transaction data ingestion to real-time pre-approval at checkout.

Why This Matters for GrabOn

GrabOn's partnership with Poonawalla Fincorp (an NBFC with full lending license) and PayU's LazyPay product makes embedded BNPL commercially viable. The 'GrabCash powered by Poonawalla Pay' concept in the PayU strategic deck is the seed of this — candidates must bring it to life as a working product prototype.

Technical Requirements

- Build a fully spec-compliant MCP server that exposes a GrabOn Transaction Data API — candidates must design the mock data schema themselves (user_id, merchant, category, GMV, coupon_used, payment_mode, return_flag, frequency)
- Implement a credit scoring engine using the MCP + Claude Code: factors must include purchase frequency, deal redemption rate (quality signal), category diversification, GMV trajectory over 12 months, and return behaviour flag
- Integrate with PayU LazyPay sandbox API for the actual BNPL disbursal flow (EMI offer generation: 3/6/9 months)
- Build a React checkout widget that shows the pre-approved EMI offer inline — must feel like a live product, not a demo
- Claude must generate a human-readable credit narrative for each user: 'You qualify because...' with specific data-backed reasoning — no black-box scores
- Add a fraud velocity check: flag if a new user tries to access BNPL within 7 days of registration

What to Submit

A working checkout demo for 5 sample user personas (ranging from 'new user with no history' to 'power user with 200+ transactions') — each getting a differentiated BNPL offer or rejection with reasoning. Must include the MCP server code, scoring logic, and checkout UI.

How It Will Be Evaluated

- MCP server is properly structured, documented, and follows the MCP specification
- Credit scoring logic is explainable and data-driven, not arbitrary thresholds
- The checkout widget is polished enough to show to a PayU/Poonawalla partner
- Claude narratives are specific, personalized, and not generic boilerplate

Project 02: GrabInsurance: Contextual Embedded Insurance at Deal Redemption

Build a system that detects purchase intent from deal category data and serves the most relevant insurance micro-product at the exact moment of redemption

[GrabInsurance] [Intent Classification] [MCP Server] [Claude API] [A/B Testing]

Difficulty	Hard — 2.5 to 3 days
Vertical	GrabInsurance (New Vertical)

The Challenge

Insurance is broken in India because it's sold, not bought — generic products pushed at the wrong moment. GrabInsurance flips this: GrabOn knows exactly what a user is about to buy (travel, electronics, health product, food) and can serve a hyper-relevant, affordable micro-insurance product in that precise context. Candidates must build the full embedded insurance flow.

Why This Matters for GrabOn

GrabOn's category mix — Travel (17%), Electronics (10%), Health (8%), Food (16%) — maps almost perfectly to insurable events. The challenge is classification accuracy and copy quality: a travel insurance offer served on a domestic flight deal must feel radically different from one served on a hotel booking deal.

Technical Requirements

- Build an intent classification MCP that takes a deal object (merchant, category, subcategory, deal_value, user_history) and outputs the top 2 insurance products to offer, with confidence scores
- Create a mock insurance product catalog with at least 8 products: Travel Cancellation, Travel Medical, Electronics Extended Warranty, Screen Damage Cover, Personal Accident (Food Delivery), Health OPD Cover, Return Journey Protection, and Purchase Protection
- Integrate with a mock insurance pricing API that returns premium quotes dynamically based on deal value and user risk tier
- Use Claude to generate hyper-personalized insurance offer copy — not 'Buy Travel Insurance' but something like 'Your ₹12,400 Goa trip. Protect it for ₹89.' — copy must be contextually accurate
- Build an A/B testing framework: for each user session, randomly assign one of 3 copy variants and track which converts (mock conversion events)
- Handle multi-category cart scenarios: user has a Myntra deal + a MakeMyTrip deal in the same session — which insurance wins, or do both get shown?

What to Submit

A working insurance recommendation engine tested against 10 mock deal scenarios across 5 categories, with A/B variant generation for each, a conversion tracking dashboard, and a sample 'insurance storefront' UI that GrabOn could embed at checkout.

How It Will Be Evaluated

- Intent classification handles edge cases (ambiguous categories, bundle deals) gracefully
- Insurance copy is genuinely contextual and personalized — evaluators will read each string
- A/B framework is properly instrumented with session IDs and event tracking
- The UI is polished enough to show to an insurance partner

Project 03: Partner Attribution & ROAS Intelligence Platform (Rakuten-Ready)

Build the 'Integrated Outcomes and Attribution Layer' that GrabOn pitches to Rakuten — a full-stack analytics platform that unifies performance across all 145 partner distribution channels

[Rakuten Partnership] [Attribution Modeling] [Claude Analytics] [ROAS/CPA] [Anomaly Detection]

Difficulty	Hard — 3 days
Vertical	Core GrabOn Infrastructure + Rakuten Strategic Partnership

The Challenge

One of GrabOn's most powerful pitches to Rakuten is: 'We can give you single-pane-of-glass visibility across 145 partner channels in India.' Right now, that is a promise. This project makes it real. Candidates must build the attribution engine, the anomaly detection layer, and the Claude-powered narrative intelligence — exactly the product a Rakuten account manager would open every Monday morning.

Why This Matters for GrabOn

Rakuten operates on ROAS and CPA metrics. GrabOn's value is that its deals are already attached to high-intent users (40M subscribers who came specifically to find deals). But without a credible attribution layer, that claim is unverifiable. This platform is the evidence layer that closes the Rakuten deal.

Technical Requirements

- Build a real-time event stream simulator that generates click, add-to-cart, and purchase events across 10 mock merchant categories at statistically realistic conversion rates (CTR ~2-4%, checkout conversion ~15-25%)
- Implement 3 attribution models side-by-side: Last Click, First Click, and Linear Position Decay — dashboard must let the user toggle between them and see how revenue attribution shifts
- Build a ROAS/CPA calculator per merchant, per channel tier (Tier 1: Amazon/Flipkart, Tier 2: mid-market, Tier 3: long-tail), and per campaign type (cashback vs coupon vs exclusive deal)
- Use Claude API to generate a structured weekly performance narrative — not a summary of numbers but actual insights: which merchant categories outperformed, which channels underdelivered, and one actionable recommendation
- Build an anomaly detection alert system: if any merchant's week-over-week conversion rate drops >20%, flag it with probable cause analysis (seasonality? competitor deal? deal expiry?)

- The dashboard must be exportable as a static HTML report that a Rakuten account manager can email to a merchant without any login

What to Submit

A fully working analytics dashboard with live event stream simulation, 3-model attribution toggle, ROAS/CPA grid, anomaly alerts panel, and a Claude-generated weekly report. Must be demoed live with the event stream running in real time.

How It Will Be Evaluated

- Attribution model logic is mathematically correct and the toggle genuinely changes revenue distribution
- Anomaly detection catches statistically significant drops with <5% false positive rate on test data
- Claude narratives read like a senior analyst wrote them — specific, data-backed, actionable
- The exportable report is clean and professional enough to send to a Rakuten stakeholder

Project 04: GrabOn × Glance: Zero-Input Deal Personalization Engine (InMobi)

Build a deal recommendation system designed for contexts where users provide zero explicit input — lock screens, push notifications, and creator storefronts — using contextual signals alone

[InMobi / Glance] [Contextual Bandits] [MCP] [Claude Recommendations] [Lock Screen UI]

Difficulty	Hard — 2.5 to 3 days
Vertical	Core GrabOn Distribution + InMobi Strategic Partnership

The Challenge

Glance's lock screen is one of the highest-reach surfaces in India — 150M+ daily active users. But it's a zero-input environment: the user hasn't typed anything, hasn't searched, hasn't expressed intent. The only signals are who they are, what time it is, what they've purchased before, and where they probably are. This project challenges candidates to build a recommendation engine that works under these constraints — and a Glance-card simulator that shows exactly how the deal would render.

Why This Matters for GrabOn

InMobi's strategic partnership with GrabOn is specifically about completing their 'Commerce Stack' — Ads (awareness) + Glance (discovery) + GrabOn (conversion). The missing piece is the intelligence layer that decides which of GrabOn's 21,000+ merchants gets surfaced at the right moment. Candidates are building that missing piece.

Technical Requirements

- Design a behavioral signal schema: time_of_day, day_of_week, location_category (home/office/transit), recent_purchase_categories (last 30 days), frequency_by_category, price_sensitivity_index (derived from coupon usage patterns)
- Build a contextual bandit recommendation algorithm (Thompson Sampling or UCB1) that ranks GrabOn's top 100 merchants for a given signal combination — no collaborative filtering shortcuts
- Use Claude to generate 160-character deal copy optimized for lock screen format: must be compelling without any contextual scaffold, not generic (not 'Great deals on electronics' but '40% off Sony headphones. 4.2★, 2,300 reviews. Tap to grab.')
- Build a Glance-card UI simulator in React: shows the card exactly as it would render on a phone lock screen, with merchant logo, deal copy, and category chip
- The system must update recommendations every 15 minutes as signals shift (simulate a user moving from 'morning commute' to 'lunch break' to 'evening at home')

- Track regret: simulate 100 recommendation cycles and show a regret curve demonstrating the bandit algorithm improves over random selection

What to Submit

A working recommendation engine with signal simulator, contextual bandit logic, Glance-card UI, and a regret curve visualization. Must demo 5 different user personas (student, working professional, homemaker, traveler, gamer) getting genuinely different deal cards at the same time of day.

How It Will Be Evaluated

- Bandit algorithm implementation is correct — not just random selection with a wrapper
- Deal copy generated by Claude is genuinely different per persona and per time context
- The Glance card UI looks like a real product, not a wireframe
- The regret curve shows measurable improvement over random baseline after 50+ cycles

Project 05: GrabCash Wallet Intelligence Platform

Build the 'GrabCash powered by Poonawalla Pay' concept — a stored-value cashback wallet with PayU checkout integration and a Claude-powered financial advisor agent

[GrabCash] [PayU Sandbox] [Poonawalla Fincorp] [Claude Financial Agent] [Wallet API]

Difficulty	Hard — 3 days
Vertical	GrabCredit + GrabInsurance Infrastructure (Wallet Layer)

The Challenge

The PayU and Poonawalla Fincorp strategic decks both converge on the same product vision: GrabCash should evolve from a passive cashback tracker into an active financial product — a stored-value wallet where users can spend on PayU merchants, transfer to bank, or invest in micro-savings products. This project is to build that product.

Why This Matters for GrabOn

GrabOn's 40M subscribers earn cashback on every transaction. Today that value is mostly unrealised — points that expire or sit in a dashboard. The GrabCash wallet turns that into a live financial instrument, deepening user stickiness and creating a new revenue stream for GrabOn via float income and investment spread on Poonawalla's NBFC platform.

Technical Requirements

- Build a mock Wallet API with full state management: balance, pending credits (cashback not yet settled), transaction history, spend categories, and a ledger of all credit/debit events
- Integrate with PayU sandbox for wallet-to-checkout spend: user selects 'Pay with GrabCash' at checkout, PayU deducts from wallet balance and confirms transaction — must use real PayU sandbox endpoints
- Build a micro-investment feature: GrabCash balance above ₹500 can be locked into a mock Fixed Deposit product via Poonawalla Fincorp's NBFC rails (mock API: 7.5% p.a. FD, 30-day minimum lock-in)
- Build a Claude-powered financial advisor agent that runs weekly: analyses the user's cashback earning pattern, spending categories, and current wallet balance, then produces a personalized recommendation — not generic advice but something like 'You've earned ₹1,840 in cashback from Zomato this month. You're likely a frequent food orderer. Consider locking ₹1,000 in GrabSave FD and keeping ₹840 liquid for the upcoming weekend.'
- Add fraud velocity detection: flag wallet transactions if balance is spent within 90 seconds of credit — alert the user and require OTP re-confirmation

- Build a wallet dashboard UI: balance card, transaction timeline, investment portfolio, and the weekly advisor recommendation card

What to Submit

A working wallet app prototype with PayU checkout integration, FD investment flow, fraud detection, and the Claude advisor agent. Must demo the full lifecycle: user earns cashback → gets advisor recommendation → invests partial amount → spends remainder at PayU checkout.

How It Will Be Evaluated

- PayU sandbox integration is real and functional — not mocked with hardcoded responses
- The advisor agent produces genuinely personalized, data-grounded recommendations
- Fraud detection logic handles edge cases (legitimate fast transactions vs anomalous ones)
- The UI is clean and demonstrates product thinking, not just engineering

Project 06: Multi-Channel Deal Distribution MCP: The 'One-to-Many' Merchant Rail

Build the MCP server that turns one merchant deal payload into six fully formatted, localized deal placements across every channel GrabOn operates — simultaneously

[MCP Server] [Claude Code] [Multi-Channel] [Localization] [Webhook Delivery]

Difficulty	Hard — 2.5 to 3 days
Vertical	GrabOn Core Infrastructure — Enables All 4 Strategic Partnerships

The Challenge

GrabOn's pitch to every partner (Rakuten, PayU, InMobi, Poonawalla) is the same: 'One integration unlocks 145 distribution pipes.' This project is to build the actual integration layer that makes that pitch true. A single merchant API call should propagate a deal across email, WhatsApp, push notification, Glance lock screen, PayU checkout banner, and Instagram — all formatted, all localized, all A/B tested.

Why This Matters for GrabOn

This is the infrastructure play that sits under every strategic partnership. When a Zomato deal is uploaded into GrabOn's system, today it takes manual work to format it for each channel. Automating this with an MCP + Claude pipeline is what allows GrabOn to scale from 3,500 to 35,000 merchant integrations without proportionally scaling headcount.

Technical Requirements

- Build a fully spec-compliant MCP server following Anthropic's Model Context Protocol specification — must be connectable to Claude Desktop as a local MCP, not just a REST API
- The MCP must expose one primary tool: distribute_deal(merchant_id, category, discount_value, discount_type, expiry_timestamp, min_order_value, max_redemptions, exclusive_flag)
- Claude generates formatted copy for all 6 channels simultaneously: Email HTML snippet (subject line + body headline + CTA), WhatsApp message (max 160 chars), Push notification (max 50 char title + 100 char body), Glance lock screen card (160 chars, must work without context), PayU checkout banner (40 chars, action-oriented), Instagram caption (with hashtags)
- Generate 3 A/B variants per channel = 18 total outputs per deal. Variants must be meaningfully different (urgency-driven vs value-driven vs social-proof-driven), not just synonym swaps

- Auto-localize all 18 outputs into English, Hindi, and Telugu = 54 total strings. Localization must be culturally accurate, not just translated (Telugu idioms for deal urgency are different from Hindi ones)
- Build a webhook delivery simulation layer: each channel has a mock endpoint that accepts the formatted deal and returns a delivery status (delivered/failed/pending). Show delivery success rates.

What to Submit

A working MCP server that candidates demo by connecting it to Claude Desktop, then asking Claude to 'distribute this deal' via natural language. The output should be all 54 formatted strings plus delivery simulation logs. Must show at least 3 merchant deals being processed end-to-end.

How It Will Be Evaluated

- The MCP server is genuinely MCP-spec-compliant and connectable to Claude Desktop — this will be tested live
- All 18 copy variants are meaningfully different from each other, not just slightly rephrased
- Hindi and Telugu localizations are culturally accurate — evaluators with native language knowledge will review these
- Webhook delivery simulation includes retry logic for failed deliveries

Project 07: AI Merchant Underwriting Agent for GrabCredit & GrabInsurance

Build an AI agent that assesses GrabOn's 3,500 active merchant partners for embedded credit limits and insurance coverage, generating explainable pre-approved offers delivered via WhatsApp

[GrabCredit] [GrabInsurance] [Claude Agent] [WhatsApp API] [Explainable AI]

Difficulty	Hard — 3 days
Vertical	GrabCredit + GrabInsurance (Merchant-Facing Products)

The Challenge

GrabCredit and GrabInsurance are not just consumer products — GrabOn's 3,500 merchant partners are equally valuable targets for embedded financial products. A small fashion merchant on GrabOn with stable GMV and growing return customers is a perfect candidate for a working capital loan or business interruption insurance. This project builds the AI underwriting agent that identifies them and makes an offer.

Why This Matters for GrabOn

Poonawalla Fincorp's NBFC license enables GrabOn to offer merchant lending. The underwriting data advantage is enormous: GrabOn has 12 months of GMV, transaction velocity, deal redemption quality, and customer return rates for every merchant on its platform. Traditional banks don't have this. This project turns that data into a working credit and insurance underwriting engine.

Technical Requirements

- Design a comprehensive merchant profile schema: merchant_id, category, monthly_gmv_12m (array), coupon_redemption_rate, unique_customer_count, customer_return_rate, avg_order_value, seasonality_index (peak vs. trough GMV ratio), deal_exclusivity_rate, return_and_refund_rate
- Build a Claude underwriting agent with two modes: GrabCredit mode (outputs working capital credit limit in ₹ lakhs, interest rate tier, tenure options) and GrabInsurance mode (outputs business interruption coverage amount, premium quote, suggested policy type)
- The agent's decisions must be fully explainable: every offer must include a 3-5 sentence rationale that references specific data points from the merchant's profile — 'We are offering ₹15L at Tier 2 rates because your GMV has grown 38% YoY, your customer return rate of 71% indicates demand stability, and your refund rate of 2.1% is below the category average of 4.8%'
- Build a risk tiering system: Tier 1 (low risk, best rates), Tier 2 (moderate risk, standard rates), Tier 3 (high risk, requires manual review) — with clear criteria for each tier

- Integrate with WhatsApp Business API (or Twilio sandbox): the agent sends a pre-approved offer message to the merchant's registered WhatsApp number. Message must be formatted as a proper business notification, not a text dump
- Build a merchant dashboard showing all 10 sample merchants: offer status, credit limit, insurance quote, risk tier, and a one-click 'Accept Offer' flow that triggers a mock NACH mandate

What to Submit

A working underwriting agent processing 10 diverse merchant profiles (covering Tier 1 through Tier 3 outcomes, different categories, including at least 2 rejection scenarios with clear explanations). Live demo of WhatsApp offer delivery for at least 2 merchants.

How It Will Be Evaluated

- Risk tiering logic is coherent and data-driven — evaluators will stress-test edge cases
- Explainability narratives are specific and reference actual data points, not template text
- WhatsApp integration is functional (sandbox is acceptable but must show real message delivery)
- The dashboard clearly shows the full underwriting decision trail for each merchant

Project 08: RankDrive 2.0: Competitive Intelligence Engine for Merchant Ops

Extend GrabOn's internal SEO tool into a full competitive intelligence platform that continuously tracks competitor coupon sites, detects pricing gaps and exclusive deals, and generates automated threat reports

[Claude Code] [Web Scraping MCP] [RankDrive Extension] [Competitive Intel] [Slack/Email Alerts]

Difficulty	Hard — 2.5 to 3 days
Vertical	GrabOn Core Operations — Merchant Retention & Deal Exclusivity

The Challenge

GrabOn already builds RankDrive internally for keyword and SEO tracking. But the merchant ops team has a bigger problem: they don't know when a competitor (CouponDunia, CashKaro, Grabon's other rivals) is offering better deals on the same merchants. By the time a merchant defects, it's too late. This project builds the early warning system — a competitive intelligence engine that monitors the landscape in real time.

Why This Matters for GrabOn

Fashion is 24% of GrabOn's GMV, Travel is 17%, Food is 16% — together they are 57% of the business. If a competitor gets exclusive deal rights on even one major merchant in any of these categories, the revenue impact is material. The merchant ops team needs intelligence, not gut feeling, to prioritize re-negotiation conversations.

Technical Requirements

- Build a crawler MCP (using Playwright or Puppeteer via Claude Code) that reads public deal pages for 10 target merchants (choose from: Amazon, Myntra, Zomato, Swiggy, MakeMyTrip, Boat, Nykaa, Puma, Ajio, CRED) across 3 competitor platforms — compare their offers to what GrabOn currently lists
- Detect and classify three types of competitive gaps: Pricing Gap (competitor offering higher discount %), Exclusivity Gap (deal only available on competitor, not GrabOn), Freshness Gap (competitor's deal updated more recently than GrabOn's by >48 hours)
- Build a merchant risk scoring system: each of the 10 merchants gets a 'defection risk score' based on the severity and recurrence of competitive gaps — flag any merchant above the risk threshold as 'Priority Re-negotiation'
- Use Claude to generate a weekly Competitive Threat Report: which merchants are most at risk, which categories are losing ground, and a specific re-negotiation strategy for each Priority merchant (what to offer: exclusivity bonus? higher cashback tier? co-funded campaign?)

- Build an automated alert system: if a competitor gains an exclusive deal in GrabOn's top 3 categories (Fashion, Travel, Food), trigger an immediate Slack message (use Slack API sandbox) AND an email alert (use SendGrid sandbox) to the merchant ops team
- The system must handle anti-scraping measures gracefully (rate limiting, User-Agent rotation) and log failed crawls for manual review

What to Submit

A working competitive intelligence dashboard with live crawl results for 10 merchants, a defection risk grid, an automated alert demo (showing a Slack message firing when a competitive gap is detected), and a PDF export of the weekly threat report generated by Claude.

How It Will Be Evaluated

- Crawler is robust and handles at least 8 of 10 merchants without failure
- Competitive gap classification is accurate — evaluators will manually verify against actual competitor sites
- Claude-generated threat report reads like a genuine strategic analysis, not a list of facts
- Alert system fires correctly and the Slack/email notifications are well-formatted and actionable

Evaluation Rubric

All submissions will be evaluated on five dimensions. Each dimension is scored 1-5 and carries equal weight in the final score.

Dimension	What We Look For	Weight
Technical Depth	MCP specification compliance, API integration quality, data model design, code architecture. Does the candidate understand what they built?	**20%**
Product Thinking	Does the submission solve a real problem at GrabOn? Is the UX considered? Would a non-technical stakeholder at Poonawalla or Rakuten find it credible?	**20%**
Demo Quality	Is the live demo polished and compelling? Does it handle edge cases gracefully? Can it survive a 10-minute Q&A?	**20%**
Claude / MCP Usage	How effectively is Claude Code used? Is the MCP well-structured? Are the AI-generated outputs (copy, narratives, recommendations) genuinely high quality?	**20%**
Code Quality & Documentation	Is the README clear? Are architecture decisions explained? Is the code readable and maintainable? Are APIs documented?	**20%**

Submission Requirements

- Public GitHub repository with clear folder structure — evaluators will clone and run your project
- A README.md that explains: what you built, why you made key architecture decisions, how to run it locally, and what you would do differently with more time
- A Loom video (10-15 minutes) walking through the live demo — must include at least one edge case or error scenario and how your system handles it
- The MCP server (for projects that require one) must be demonstrably connectable to Claude Desktop — we will test this live during evaluation
- If using paid API sandboxes (PayU, WhatsApp, SendGrid), include either sandbox credentials or clear mock fallbacks that demonstrate the integration architecture

Good luck. We are looking for builders who make us think — 'We should hire this person immediately.' That is the bar.